

# Conner Jordan

San Luis Obispo, CA

📞 (805) 975-9793 — 📩 [connercharlesjordan@gmail.com](mailto:connercharlesjordan@gmail.com) — 📍 San Luis Obispo, CA  
LinkedIn: [linkedin.com/in/conner-jordan-4b268514a](https://linkedin.com/in/conner-jordan-4b268514a) — GitHub: [github.com/cjordan223](https://github.com/cjordan223)

February 24, 2026

Hiring Manager  
fastino.ai

Dear Hiring Team,

I am applying for the AI Platform Engineer role at fastino.ai because I build secure, production-grade infrastructure where data pipelines meet operational reliability. My experience architecting distributed ingestion systems and normalizing inconsistent security data aligns directly with your need for scalable fine-tuning pipelines, especially where messy real-world inputs must feed structured training workflows.

At the University of California Office of the President, I designed and deployed containerized services that ingest multi-source security telemetry, reconcile conflicting asset records, and trigger remediation workflows at scale. I implemented runbooks and guardrails that made these pipelines auditable and repeatable, which mirrors your emphasis on experiment tracking and reproducibility. I also built AI-assisted tooling using LangChain and RAG components to surface actionable insights from security logs, demonstrating hands-on experience integrating model inference into operational systems.

Earlier, I developed Python and PowerShell automation for enterprise endpoint management, including configuration enforcement and patching pipelines, that required tight integration with cloud infrastructure and version-controlled workflows. This background gives me a strong foundation for designing CI/CD-oriented experiment tracking systems and dataset versioning pipelines. My phishing detection project combined a browser extension with a backend API and earned an innovation award, illustrating how I build end-to-end systems where model behavior is observable, testable, and maintainable.

I ship systems that endure beyond the first deployment: they are documented, monitored, and hardened against real-world variability. At fastino.ai, I would bring that same operational discipline to your AI infrastructure, ensuring LoRA and distillation workflows are reproducible, GPU utilization is optimized for cost and throughput, and reward modeling loops remain traceable and secure. I thrive where reliability, security, and developer velocity intersect.

Thank you for your time and consideration. I would welcome the opportunity to discuss how my background in secure, scalable infrastructure can support fastino.ai's mission of building reliable AI systems that work in production.

Sincerely,  
Conner Jordan